

# BG2 Pod: The SpaceX IPO, Fable 5 & the AI Capex Reality Check

Extracted from: BG2 Pod (Brad Gerstner) — w/ Gavin Baker, Andrew Fox & Clark Tang

Source: [BG2 Pod \(YouTube\)](#)

Published: 11-06-2026 · 1h 20m · transcript archived at `transcripts/20260611_bg2_spacex_ipo_fable5_ai_capex.{json,txt}`



**7-Minute Version:** [Watch Key Excerpts](#)

---

## 0) Who Is On It — And Who Does What

| Person                     | Firm / Role                                                                                                      | Role in this episode                                                                                                                                                    | Skin in the game (researched, sourced)                                                                                                                                                                                                                                                                                                                                                                                                                      |
|----------------------------|------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Brad Gerstner</b>       | Founder/CEO, Altimeter Capital (~\$10–20B AUM; ~29.5%/yr public book since 2011)                                 | Host. Frames every segment, supplies the bank-model numbers (\$160B 2028 rev, \$1.5T capex vs \$300B inference revenue), runs the “market check” close                  | Pre-IPO SpaceX holder — but <b>publicly stated he’s SELLING ~30% into the IPO</b> (fiduciary duty, All-In Liquidity Summit) while saying “must buy, must own” on air; <b>led Anthropic’s \$65B Series H at \$965B post (May 28, 2026)</b> ; OpenAI investor; NVDA = top 13F position (~\$1.5B, ~8.1M sh); doubled CoreWeave to ~\$325M in Q4’25; Cerebras investor; drove the Trump-accounts/ Invest America initiative (NVDA among corporate contributors) |
| <b>Gavin Baker</b>         | Founder/CIO, Atreides Management (~\$4B; ex-Fidelity OTC PM, 19.3%/yr, top-1% of peers)                          | The intellectual engine. Pareto-curve framework, Cursor/data thesis, Einstein long-running-agent argument, Nvidia-vs-ASIC strategy, market-runner analogy               | <b>SpaceX = Atreides’ largest private position</b> ; he made Fidelity’s first SpaceX investment in 2015 at ~\$10B valuation (~175x to IPO price); originated the public “Elon Web Services” thesis; bought <b>leveraged NVDA call options Feb 2026</b> ; advisory-level intimacy with xAI (“I asked the team... be a little less public about data-center design”); to his credit, routinely discloses “I am biased given my Atreides investments”          |
| <b>Andrew Fox (“Foxy”)</b> | Atreides Management (analyst, joined Jan 2021; ex-PJT Partners TMT; ASU)                                         | The quant of the table. Launch economics, \$/kg math, orbital-vs-terrestrial \$/GW capex, monetization-per-GW trajectory, training-vs-inference capex split             | Same book as Baker. Near-zero public footprint — this episode is one of his most public appearances                                                                                                                                                                                                                                                                                                                                                         |
| <b>Clark Tang</b>          | Partner, Altimeter (youngest ever; ex-Goldman/CS, Valor Equity — Anduril/Reddit; II 2025 Hedge Fund Rising Star) | Asia/semis channel-checker. Fresh from Computex/GTC Taiwan. Author of the xAI operating-profit-per-gigawatt chart; Vera Rubin allocation intel; MediaTek/ASIC landscape | Altimeter positions; <b>pitched NVDA internally Sept 2022</b> (the firm’s defining trade); his charts are the analytical backbone of the SpaceX bull case being marketed                                                                                                                                                                                                                                                                                    |

**Notably absent: Bill Gurley** — the show’s co-host and resident skeptic, who argued the open-source/asymptote side of this exact debate “at this table two years ago.” The bear chair was empty by construction.

**Bottom line on the table:** four people, two funds, all long SpaceX into the IPO (one of them simultaneously *selling* into it), Altimeter holding stakes in **three of the “next \$1T revenue” trio it forecasts** (SpaceX, Anthropic-lead-investor, OpenAI), plus NVDA/CoreWeave/Cerebras across both books. This is the most conflicted panel BG2 has run — *and they say so*: “none of us are here to pump the IPO” (immediately after Gerstner calls it “a must buy, a must own... set it and forget it”). Full bias audit in §6.

## Executive Summary

The episode is a **pre-IPO bull case for SpaceX delivered by its own shareholders on pricing day** (\$135/share, \$1.77T, Nasdaq:SPCX, first trade the next morning — with a record ~30% retail allocation), wrapped in genuinely

high-quality analysis. The single most important external fact: **Gerstner separately stated he's selling ~30% of Altimeter's SpaceX into this IPO** while calling it “a must buy, must own” on air. The frame that matters: SpaceX is no longer a launch + Starlink story but a **four-lever compound** — (1) launch/rapid reusability, (2) Starlink broadband + direct-to-cell, (3) a brand-new AI-compute resale business (“Elon Web Services”) that went from non-existent to **#4 hyperscaler in 30 days** via the Anthropic (\$22–23B/GW) and Google (\$50B/GW) deals, and (4) the xAI frontier model supercharged by the **Cursor acquisition** — which Gerstner explicitly flags as “the least talked about, biggest upside surprise.” The deals added \$29B of revenue in a month and compressed the IPO multiple from ~100x to ~39x trailing — a number that does the heavy lifting in the valuation defense.

The second act is the more durable alpha. **Long-running agents broke the benchmark regime:** Fable 5 / ChatGPT 5.5-class models mean “we no longer know how smart models are” because nobody can evaluate a model for a year before the next one ships (the Noam Brown point). Baker's conclusion — “however bullish I was on compute before, I'm a lot more bullish” — feeds directly into the episode's resolution of the open-source debate: **frontier captures ~90% of revenue while open source carries ~80% of tokens, and both persist** — which is *bullish for compute/hardware regardless of which layer wins margin*. On capex math: \$1.5T 2027 capex vs ~\$300B inference revenue “maths” at 60–70% gross margin, monetization-per-GW *rising* (\$20B → \$30–40B/GW within a year), and <0.2% of humanity using AI agentically.

The closing market check is where the smart money diverges from the pitch: **both managers have already de-risked** (Altimeter “large → medium-small”), both expect a consolidation in a seasonally weak summer (token consumption has plateaued each of the last three summers), and Baker warns the “find the next bottleneck” trade “was the last game — that game is over” — directly validated by our data (CEG –30% YTD, VRT –19%/1M, NVDA/AVGO lagging SMH by ~58pp YTD). The honest read: structurally maximal bulls, tactically cautious, telling retail “must own” on the same tape they trimmed.

## Listen or Skip?

**Verdict: Must-listen** — with the bias filter on at all times.

| Aspect                     | Rating |
|----------------------------|--------|
| Signal density             | ★★★★★  |
| Actionability              | ★★★★☆  |
| Novel insights             | ★★★★☆  |
| Independence / objectivity | ★★☆☆☆  |

**Best for:** anyone underwriting the SpaceX IPO, sizing AI-compute exposure, or wanting the sharpest articulation of the frontier-vs-open-source economics. **Skip if:** you want a bear case — there is no bear in the room; the closest is their own positioning behavior.

## 1) Episode Map

| Time    | Segment                                                              |
|---------|----------------------------------------------------------------------|
| 00:40   | Intro — SpaceX IPO in 2 days, Mythos launch, Taiwan takeaways        |
| 03:05   | xAI's Google & Anthropic deals — highest operating profit per GW     |
| 13:28   | "Elon Web Services" — AI compute wasn't in anyone's SpaceX model     |
| 19:08   | Data centers are not commodities — first-principles design, 122 days |
| 26:01   | Orbital compute economics — \$5B/GW space vs \$20–25B/GW ground      |
| 29:25   | Cursor — the most underrated variable for xAI's model                |
| 35:05   | Bull & bear — can SpaceX 8x revenue in 4 years?                      |
| 37:00   | Post-IPO drawdowns, lock-ups, sizing                                 |
| 43:56   | Fable 5, Mythos & why snapshot benchmarks are broken                 |
| 51:00   | Frontier vs open source — 90% of revenue accrues at the frontier     |
| 1:04:11 | \$1.5T capex vs \$300B AI revenue — does the math math?              |
| 1:18:13 | The next \$1T of revenue: three companies, half the time             |

## 2) Core Theses (with attribution)

### T1 — SpaceX is a four-lever compound; two variables dominate (*Baker*)

**Claim:** The IPO hinges on (a) how fast terrestrial data centers come online and (b) what Cursor's data does to xAI's model. Everything else (launch, Starlink) is "doable." **Mechanism:** Colossus 1 at 55% IRR (Freda/Altimeter calc) vs 6–8% cost of debt; 122-day builds vs industry 4 years; monetizing at \$22–50B/GW vs ~\$14B/GW implied by the banks' \$160B 2028 number — i.e., **the leaked bank models already look conservative on rate, so volume (GW online) is the swing factor**. **Confidence as stated:** High on terrestrial compute, high on Starlink scale-up, optionality on orbital + model. **Our read:** the decomposition is genuinely good analysis. But note every "lever" is sourced from the same in-group (Clark's chart, Freda's IRR, Jensen's quotes) — no external check on the \$50B/GW Google rate's durability, and the 90-day termination clauses got one passing sentence.

### T2 — Rapid reusability is the master variable everyone should watch (*Fox*)

**Claim:** Two-stage rapid reuse (Starship flying 30–50x before retrofit) takes launch from ~\$1,500/kg (Falcon) to ≤\$250/kg, asymptoting to fuel cost. This unlocks Starlink V3, direct-to-cell, AND orbital compute. **Falsifiable milestones given:** second-stage recovery attempt **later this year (2026)**, re-fly **next year (2027)**, cadence from ~165/yr → "high hundreds" → thousands within ~3 years thereafter. **Our read:** this is the cleanest watch-item in the episode — binary, dated, public. If second-stage recovery slips into 2027, the orbital math and Starlink V3 numbers all shift right.

### T3 — Orbital compute: \$5B/GW (space) vs \$20–25B/GW (ground) on the non-silicon half (*Fox + Clark*)

**Claim:** 100t per Starship ≈ 5MW of satellite compute per launch → ~\$5B/GW to orbit vs \$20–25B/GW for terrestrial land/shell/power/cooling. All-in: ~\$30B/GW orbital vs ~\$60B/GW terrestrial — with the ground number *inflationary* and the space number *deflationary*. **Key honesty point:** both Fox and Baker stress orbital is **NOT needed for the IPO valuation** — it's a call option. Reliability/maintenance (GPUs melt, lasers fail) is the conceded unknown. **Our read:** the BOM math is coherent but assumes T2 fully succeeds; treat as a 2028+ option, worth ~nothing in a 2026 underwriting and everything in the bull tail.

**T4 — Cursor is the most underrated variable (*Gerstner's pick, Baker's mechanism*)**

**Claim:** Cursor + Anthropic each hold more proprietary coding tokens than exist on the public internet (*UNVERIFIED* — *Baker hedges "my understanding"*). Composer 2.5 (Kimi K2.5 base + Cursor data + 3 weeks on Colossus 2) went Pareto-dominant on coding. Now the same data is being injected into **Grok 4.3 pre-training (1.5T params)** — if scaling laws hold, xAI jumps to a frontier coding model with massive owned compute behind it. **Why it matters:** coding is "bitter-lesson adjacent" — the fastest path to AGI (Amjad Masad cite), and the highest-monetizing workload (Anthropic existence proof). **Caveat the panel skipped:** external reporting says Cursor is a **\$60B right-to-acquire later in 2026** (with a \$10B collaboration fallback), not a closed deal — the panel speaks of it as completed. The "downloaded the team into SpaceX" framing is ahead of the paperwork. **Watch:** the first Grok 4.3+Cursor-data release is the single highest-information event for the xAI line item — plus whether the \$60B option is actually exercised.

**T5 — Long-running agents broke benchmarks; nobody knows how smart models are (*Baker, via Noam Brown*)**

**Claim:** Fable 5 / ChatGPT 5.5-class models do long-horizon tasks (Stripe: 50M-line Ruby refactor in a day). Snapshot benchmarks are obsolete — the x-axis must be time/tokens. "Nobody has run Mythos for a year continuously... we may never know how smart each generation actually is." **Implication drawn:** compute demand is *understated* — you have to "let the horses run" for long periods to extract the intelligence. Einstein-for-a-year thought experiment. **Our read:** this is the genuinely new macro input of the episode. If model evaluation now requires wall-clock time × compute, then inference demand scales with *agent runtime*, not query count — structurally different demand curve than 2024-25 chatbot economics.

**T6 — Frontier gets ~90% of revenue; open source gets ~80% of tokens; BOTH persist (*Baker, Clark dissents mildly*)**

**Claim:** Every frontier release unlocks new use-case revenue (this year: long-running agents). Enterprises route commodity work to open source (Harvey: OSS + own data + router beat Opus 4.7/4.8 at lower cost) but pay frontier rates for high-value work. **Open source winning share is BULLISH for compute/hardware** — less margin at the model layer = more spend at the silicon layer. **Clark's dissent (worth weighting):** Asia consensus is the opposite of Palo Alto consensus ("right model to right workload, don't overspend"); *next 12 months decide it*; a Silicon-data-index shift toward cheap open-source tokens already appeared in the last 2 weeks. **Our read:** the 90/80 frame is the most portable mental model in the episode. Clark's seasonality+routing caveat is the early-warning indicator: watch token-mix data, not benchmark releases.

**T7 — Nvidia held share; the ASIC threat is real but two-sided (*Clark + Baker*)**

**Claim:** Everyone predicted Nvidia share loss for years; it didn't happen (ex-Anthropic, they arguably *gained*). OpenAI's 27GW paper allocation (NVDA 10 / AVGO 10 / AMD 6+warrants / Cerebras 1) implies ~37% NVDA share — Baker: "I'll be very surprised if that is where they land" (he means NVDA lands *higher*) because **tokens-per-watt = revenue in a watt-constrained world**. Meta & Microsoft ASICs disappointing; OpenAI's "Jalapeno" chip good but runs hot→cooling cost. New vector: MediaTek V8T vs Broadcom for custom workload-specific ASICs (Computex takeaway). **The strategic kicker (Baker):** Jensen can weaponize open source — "That's a cute ASIC... how would you like open source to join the frontier?" If NVDA ships frontier-grade open models, it guts the margin pool that funds customer ASIC programs, and NVDA "can become one of the world's largest cloud computing companies much faster than people think." **Our read:** the most actionable public-market section. NVDA/AVGO are the *laggards* of the complex YTD (verified below) while the bottleneck names already ran — the table is implicitly telling you where they think the next rotation goes.

**T8 — The capex math maths — for now (*Gerstner frames, all endorse*)**

**Claim:** 2027 capex ≈ \$1.5T (MS at \$1.1T + SpaceX/CoreWeave/etc.) vs ~\$300B 2027 inference revenue. Defenses: (a) 60–70% GM on that revenue, (b) ~35% of capex is training (non-revenue-bearing R&D, falling as a share), (c) revenue trajectory \$200B exit-2026 → \$400–500B 2027 → "trillions before 2030" (Dario), (d) monetization per GW rose from ~\$20B to \$30–40B in six months ("accidental profitability"), (e) <0.2% of Earth uses AI agentically. **Anti-thesis trigger (Gerstner's own):** "the second we get nervous about [the spend], the entire semi complex comes down a lot." **Our read:** the bull math requires *both* the revenue ramp and pricing holding. Clark's open-source-token-shift observation is the leading indicator that pricing may not hold across the whole curve.

**T9 — Market check: trimmed, expecting consolidation, structurally higher (*Gerstner + Baker*)**

**Stated positioning:** Altimeter went **large** → **medium-small** after the April–May run. Baker: stocks "went straight up a cliff... they're tired, they need to rest"; "finding the next bottleneck was the last game — that game is over."

Seasonal AI-summer slowdown (3rd year running), CPI back above 4 (4.2 headline, core 0.2 m/m better than 0.3 exp), oil ~\$100, Iran war risk. Both end structurally bullish: hard to be bearish post-Fable-5; next \$1T of revenue from 3 companies in half the time of the Mag 7's second trillion (\$1T rev → \$17T mkt cap precedent). **Our read:** THE tell of the episode. Pitch = "must own, set and forget." Book = trimmed to medium-small, waiting for the consolidation. Both can be honest (horizon difference), but retail listeners will hear only the first half.

### 3) Hard Data vs Narrative — Verification Layer

#### Claims verified against our local data (bars through 2026-06-11)

| Podcast claim                                     | Our data                                                                              | Verdict                                                          |
|---------------------------------------------------|---------------------------------------------------------------------------------------|------------------------------------------------------------------|
| "Semis ripped... doubles and triples"             | SMH <b>+69.2% YTD</b> ; MU +248.9%, MRVL +230.6%, AMD +128.1%                         | Confirmed                                                        |
| "Nvidia and Broadcom have been laggards" (Baker)  | NVDA <b>+9.9% YTD</b> (–13% off high), AVGO <b>+11.4%</b> (–20% off high) vs SMH +69% | Confirmed — striking dispersion                                  |
| "Software down 8% on the year"                    | IGV <b>–14.0% YTD</b> (CRM –37%, NOW –33%)                                            | Direction right, magnitude worse on our basket                   |
| "Internet down 16%"                               | FDN –1.5% YTD, but MSFT –19%, META –14%                                               | Basket-dependent; mega-cap internet names confirm, index doesn't |
| "SPY/NASDAQ up but only because of AI components" | SPY +8.5%, QQQ +16.7%, equal-weight RSP +~3%                                          | Confirmed (breadth narrow)                                       |
| "Last week some retracement"                      | NVDA –7.2%/1M, AVGO –8.0%, GOOGL –7.6%, AMZN –9.1%                                    | Confirmed                                                        |
| "Next-bottleneck game is over"                    | CEG –30% YTD, VST –9%, VRT –19%/1M, GEV –15%/1M                                       | Power/cooling trade already broken                               |
| CPI 4.2 headline / core 0.2 vs 0.3 exp            | (stated on recording morning)                                                         | Data — consistent with our macro context                         |

#### Claims verified externally (web research, sourced in companion brief)

| Podcast claim                                             | External record                                                                                                                                                                                                                                   | Verdict                                                                                              |
|-----------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------|
| IPO \$135/share, \$1.77T, "in two days"                   | Confirmed: priced June 11 after close; <b>first trade June 12</b> (next morning, not 2 days); Nasdaq ticker <b>SPCX</b> ; 555.5M shares; <b>\$75B raised — largest IPO in history</b> ; ~30% retail allocation; Musk >82% voting control post-IPO | Terms confirmed; episode dropped at pricing                                                          |
| Anthropic + Google deals "added \$29B revenue in a month" | S-1/press: Anthropic <b>\$1.25B/mo through May 2029</b> (Colossus 1, ~220K GPUs, 300MW ≈ \$15B/yr); Google <b>\$920M/mo Oct 2026–Jun 2029</b> (~\$11B/yr). Combined ≈ <b>\$26B annualized</b>                                                     | Directionally right; \$29B modestly overstated vs \$26B disclosed                                    |
| "We bought Cursor" (discussed as completed)               | Reported as a <b>\$60B right-to-acquire later in 2026</b> , with a \$10B collaboration alternative (announced April 2026)                                                                                                                         | <b>Material nuance omitted on air</b> — the model/data upside is an option, not a closed acquisition |
| 122-day data center builds                                | Reported sequence: <b>122 → 91 → 66 days</b> per facility in 2026 iterations                                                                                                                                                                      | Confirmed, and the on-air number is stale-conservative                                               |
| xAI merger context                                        | SpaceX acquired xAI <b>Feb 2026</b> (~\$1.25T combined; xAI ~\$80B); S-1 (May 20) shows a <b>\$4.94B loss absorbed</b> from the merger — never mentioned on air                                                                                   | The merger loss is the kind of number the episode skips                                              |
| Elon ~50% owner, 365-day lockup                           | Musk holds <b>&gt;82% voting control</b> post-IPO (economic ~50%ish consistent)                                                                                                                                                                   | Consistent                                                                                           |

## Headline numbers from the episode (full inventory in §8)

- **IPO:** \$135/share, \$1.77T, ~2 days post-recording. ~100x TTM revenue pre-deals → ~39x after +\$29B/month of signed deal revenue. 2025 revenue ~\$18B; banks model **\$160B by 2028** (≈8x in 3-4 yrs).
- **Deal economics:** Anthropic \$22–23B/GW, Google **\$50B/GW** (premium read partly as a “call option on first-in-line for space compute”); implied bank-model monetization only \$14B/GW → bank numbers conservative on rate.
- **Orbital:** \$5B/GW (space, non-silicon) vs \$20–25B/GW (ground); all-in ~\$30B vs ~\$60B/GW; 5MW per Starship launch; \$1,500/kg → \$250/kg → fuel-cost asymptote.
- **Compute landscape:** xAI ≈ **20% of early Vera Rubin capacity** (Clark channel checks, UNVERIFIED); OpenAI paper allocations NVDA 10GW / AVGO 10GW / AMD 6GW / Cerebras 1GW; world adds 20–25GW of AI DC capacity in 2026.
- **Capex vs revenue:** ~\$1.5T 2027 capex vs ~\$300B inference revenue; monetization per GW \$20B→\$30–40B in ~6 months; Mag-7 FCF down ~80%; the \$250B of private raises = “1% of Mag 7 market cap.”
- **Lock-ups:** Elon ~50% holder, locked 365 days; dribble-out release schedule until first earnings; BUT employees/investors had ~20 liquidity windows over 10 years (less pent-up selling than typical IPO); average max post-IPO drawdown of 20 mega-IPOs >50%.

## 4) Key Quotes

1. *“If you can borrow money at 6, 7, 8% and invest in something with a 55% IRR — I’m not the most sophisticated thinker, but that math maths.” — Baker [0:03:27]*
2. *“In 30 days, we went from not being an AI hyperscaler to being number four. And we passed a lot of companies, including Oracle.” — Baker [0:18:27]*
3. *“We do not know how smart these models are... we may never know how smart each generation of models actually is, because we don’t have time to appropriately evaluate their intelligence before the next model comes out.” — Baker [0:44:38]*
4. *“Imagine Albert Einstein had just thought about fundamental physics 24 hours a day... and he thought for one year. We might already have solved a lot of these intractable problems.” — Baker [0:45:44]*
5. *“Frontier might be 90% of the economic value. Open source might be 80% of tokens... the better open source does, the better it is for compute providers.” — Baker [0:53:28–0:54:02]*
6. *“That’s a cute ASIC you’ve built there. That is so cute. How would you like open source to join the frontier? How do you like them apples?” — Baker, ventriloquizing Jensen [0:57:26]*
7. *“Less than 0.2% of people on Earth are actually using AI in an agentic way.” — Gerstner, citing Alex/Whale Rock [1:07:58]*
8. *“Everyone expected token pricing to be smoothly deflationary. This year we’ve seen the opposite... monetization rates per watt are increasing.” — Fox [1:07:22]*
9. *“Finding the next bottleneck — I think that was the last game. That game is over... A lot of these stocks have gone straight up a cliff. They’re tired. They need to rest.” — Baker [1:14:38]*
10. *“In the last seven years we added \$1 trillion of revenue to the Mag 7, and that added \$17 trillion in market cap. The forecast now is another trillion of revenue from just three companies — SpaceX, Anthropic, OpenAI — in half the time.” — Gerstner [1:17:57]*
11. *“It’s a must buy, a must own. It’s set it and forget it.” — Gerstner, on SpaceX [0:00:00 / 0:36:19]*
12. *“I always assume a bullet is coming for me. Head on a swivel. It’s the bullet you don’t see that gets you.” — Baker [1:16:51]*



## 5) Where They're Most Likely Wrong (anti-theses)

1. **Deal-rate durability.** The \$50B/GW Google rate is treated as a monetization benchmark; it's more plausibly a scarcity-clearing price + space-option premium that mean-reverts as 2026-27 capacity lands. The 90-day termination clauses got one dismissive sentence. If Google re-rates to Anthropic-deal levels, the "39x with rising rates" framing decays fast.
2. **The 8x revenue base rate.** Gerstner names it then talks past it: almost no company has 8x'd from an \$18B base in 3-4 years. Every lever has to hit simultaneously — and T2's own milestones (second-stage reuse not yet attempted) sit upstream of two of the four levers.
3. **Circularity in the inference-revenue defense.** The \$300B revenue that justifies \$1.5T capex includes labs' revenue from selling compute/API to each other and to compute-funded startups. Nobody on the table raised vendor-financing or circular-revenue risk (AMD warrants to OpenAI being a mild form).
4. **Token-pricing inflation as a permanent regime.** Fox's "monetization per GW rising" is six months of data in a shortage. Clark's own observation — open-source token share shifting in the last 2 weeks + 3rd consecutive seasonal summer plateau — is the counter-evidence sitting inside the same episode.
5. **Selection bias on the panel.** Four bulls, two firms, zero bears, all long the asset being IPO'd in 48 hours. Bill Gurley — the show's resident skeptic, who argued the open-source/asymptote side "at this table two years ago" — is conspicuously absent from this episode.

## 6) BS & Bias Filter — Who's Talking Their Book

Everyone, with varying degrees of disclosure. Ranked by conflict intensity:

- **Gerstner (highest — and now documented):** Pre-IPO SpaceX holder who tells listeners it's "a must buy, a must own, set it and forget it" — while having **publicly stated (All-In Liquidity Summit) that he is selling ~30% of Altimeter's position into this IPO** as fiduciary duty. On the pod he says "we also expect to be buying in the IPO" — both can be technically true (sell pre-IPO secondaries, buy IPO allocation, net reduce), but the asymmetry between what retail hears and what his fund does is the single most important fact about this episode. The conflict web goes further: Altimeter **led Anthropic's \$65B Series H two weeks before this episode** — so when he credits "Anthropic showing up with revenue" for saving the 2026 tape and forecasts "the next \$1T of revenue from SpaceX, Anthropic, OpenAI," he holds stakes in **all three**. NVDA is his top public position (~\$1.5B) and NVIDIA is a corporate contributor to his Trump-accounts initiative; CoreWeave was his biggest Q4 add (doubled to ~\$325M). Counterweight to credit: he shows the >50% post-IPO drawdown chart himself and admits Altimeter cut gross exposure to "medium-small."
- **Baker (high, but the most epistemically honest):** SpaceX is **Atreides' largest private position** and his personal 11-year compounder (first check at ~\$10B valuation in 2015 → \$1.77T IPO ≈ 175x). He originated the "Elon Web Services" framing this episode mainstreams, and bought leveraged NVDA calls in February — every theme he advances is in his book. BUT his history is of disclosing exactly that ("I am biased given my Atreides investments"), and on the pod he inserts the hedges himself: rapid reusability "is a really hard thing," "I don't know what's going to happen in the short term," orbital not needed for valuation, refuses to forecast the two key variables. Talking his book, yes — while flagging the load-bearing assumptions.
- **Clark Tang (medium):** his operating-profit-per-GW chart and "20% of Vera Rubin" channel checks are the analytical pillars of the bull case, produced inside the firm that owns the asset — and NVDA was his career-making internal pitch (Sept 2022). He's also the only one voicing the Asia-consensus counterview on open source and the token-mix warning — the closest thing to a bear in the room.
- **Andrew Fox (medium):** the math is presented straight and the assumptions are explicit (\$/kg, MW/launch, BOM splits). Same book as Baker, but his sections are the most replicable-from-first-principles.
- **Structural bias:** the episode functions as **pre-IPO marketing with high production values**, released the day the IPO priced (\$135, June 11) with first trade the next morning — and with a record ~30% retail allocation to absorb. The "quasi-public, 20 liquidity windows, less selling pressure" argument and the "dribble-out lockup" chart are exactly what a banker would want circulating at pricing. None of this makes the analysis wrong — Colossus IRR, deal rates, and the Cursor logic stand on their own — but every soft judgment ("must own") should be discounted hard.

**Track-record context:** Gerstner/Baker were loudly right on AI compute 2023-2026 (Gerstner bought NVDA Q4'22 pre-ChatGPT; Tang pitched it; Baker's Fidelity OTC record is top-1%) and the panel correctly notes the consensus they violated (open-source catch-up, token deflation). Their hit rate on *this* theme has been excellent — which is exactly when book-talking is most persuasive and least checked.



## 7) Companies & Sectors Map — How To Think About Each

### A. The IPO complex (private → public)

| Name                                                                | What was said                                                                                                                                                                                                                                                                                        | How to think about it                                                                                                                                                                                                                                                                                                                                                                        |
|---------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>SpaceX / xAI</b> (IPO 2026-06-13, \$135, \$1.77T)                | Four-lever compound; “must own” per Gerstner; 39x trailing post-deals; Elon locked 365d                                                                                                                                                                                                              | The analysis is strong, the messengers are conflicted. If participating: expect >50% drawdown risk per their own chart; their own playbook (base position + ballast, size up into post-IPO weakness) is better advice than “set and forget.” Highest-information catalysts: Starship 2nd-stage recovery attempt (H2’26), first Grok 4.3+Cursor release, first earnings (lockup dribble ends) |
| <b>Cursor</b>                                                       | “Most underrated variable”; ~\$10B exit-2026 revenue (Altimeter est.); proprietary coding tokens > public internet (UNVERIFIED). <b>Correction: not yet acquired</b> — SpaceX holds a \$60B right-to-acquire later in 2026 (vs \$10B collaboration alternative); the panel discussed it as if closed | The embedded call option in the SpaceX IPO — but an <i>unexercised</i> one. Watch (a) whether the \$60B option is exercised, (b) Composer 3 / Grok 4.3-coding benchmarks as the option going in/out of the money                                                                                                                                                                             |
| <b>Anthropic</b>                                                    | Revenue “showed up in a major way” — credited with saving the whole 2026 tape; Fable 5/Mythos = frontier leadership; paying SpaceX \$22-23B/GW                                                                                                                                                       | The de facto index of AI inference revenue. Their disclosed revenue trajectory is THE macro datapoint for the entire semis complex — more important than any Fed print for this tape                                                                                                                                                                                                         |
| <b>OpenAI</b>                                                       | ChatGPT 5.5 long-running capabilities; 27GW paper allocations; “Jalapeno” ASIC good but hot                                                                                                                                                                                                          | Watch the gap between paper GW commitments and deployed GW — Baker expects NVDA’s realized share to beat the 37% paper share                                                                                                                                                                                                                                                                 |
| <b>CoreWeave + NeoClouds</b> (Nebius, Iren, “50 more being funded”) | “Huge business, we’re investors”; but SpaceX just became the #4 hyperscaler overnight and may consolidate the category                                                                                                                                                                               | The uncomfortable subtext: EWS is a <i>competitor</i> to every NeoCloud. If SpaceX can stand up GW faster/cheaper and borrow cheaper, second-tier NeoClouds get squeezed on both cost of capital and customer access. Be selective; the category’s “shortage rent” is what EWS arbitrages away                                                                                               |
| <b>Cerebras</b>                                                     | 1GW of the OpenAI allocation; both funds invested; Cerebras IPO’s dribble-out lockup cited as precedent                                                                                                                                                                                              | Real allocation win, but 1/27 of the paper pie; same caveat as all ASICs — tokens-per-watt vs NVDA decides economics                                                                                                                                                                                                                                                                         |
| <b>Reflection</b> (US open-source lab)                              | Baker “very impressed” but neither fund invested; wants a US frontier OSS lab to win                                                                                                                                                                                                                 | The tell is that neither fund put money in despite wanting it to exist — revealed preference says OSS labs capture tokens, not value                                                                                                                                                                                                                                                         |

## B. Semis & hardware (public)

| Name               | What was said                                                                                                                        | Our data overlay (RS pct 27d / long-horizon, 2026-06-11)                 | How to think about it                                                                                                                                                                                                                                 |
|--------------------|--------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>NVDA</b>        | Held/gained share vs every prediction; tokens-per-watt = revenue; can weaponize open source + become a top cloud; Vera Rubin ramping | 58th pct / 76th — mid-pack, –13% off high, +9.9% YTD                     | The episode's implicit rotation candidate: fundamentals-positive narrative + laggard price. Baker expects realized OpenAI share > paper 37%. The open-source-frontier option is unpriced. Strongest risk-reward in the complex per this table's logic |
| <b>AVGO</b>        | 10GW OpenAI allocation; but MediaTek now attacking custom-ASIC space (V8T vs Broadcom)                                               | 22nd pct / 89th — classic Pullback-Buy profile (252d leader, 27d washed) | Long-horizon leadership intact, swing-weak. The MediaTek threat is the new variable — watch Computex follow-through. Conviction taxonomy says P (pullback buy) not breakdown                                                                          |
| <b>AMD</b>         | 6GW + warrants from OpenAI                                                                                                           | 98th pct / 98th — top decile both horizons, +128% YTD                    | Already paid. The warrants = circular-revenue flag. Momentum says hold, valuation work needed before adding                                                                                                                                           |
| <b>MU (memory)</b> | Memory repeatedly cited as Taiwan bottleneck topic                                                                                   | 99.8th pct / 99.4th — #1 in our universe, +249% YTD                      | The “found bottleneck” that already 3.5x'd. Baker's cliff-climber warning applies most literally here. Late-cycle momentum; trail stops, don't initiate                                                                                               |
| <b>MRVL</b>        | (implied in ASIC/custom-silicon discussion)                                                                                          | 99.7th / 96th, +231% YTD                                                 | Same as MU — the bottleneck game's biggest winners are the most “tired runners”                                                                                                                                                                       |
| <b>TSM</b>         | Computex backdrop; everything routes through them                                                                                    | 63rd / 92nd, +38.9% YTD, –5.7% off high                                  | The boring structural hold of the whole complex; least narrative risk per dollar of AI exposure                                                                                                                                                       |
| <b>MediaTek</b>    | NEW: V8T custom ASIC vs Broadcom — Computex's biggest topic per Clark                                                                | (not in universe)                                                        | The new entrant in custom AI silicon. Worth adding to monitoring universe (TW-listed; US ADR illiquid)                                                                                                                                                |
| <b>Cerebras</b>    | see above                                                                                                                            | —                                                                        | —                                                                                                                                                                                                                                                     |

### C. Hyperscalers & mega-cap tech

| Name               | What was said                                                                                                                             | Our data overlay                                                    | How to think about it                                                                                                                                                                          |
|--------------------|-------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>GOOGL</b>       | Paying SpaceX \$50B/GW premium (partly space-compute option); raised \$80B; Gemini 3.1 Pro on the frontier; TPU+Anthropic model validated | 24th pct / 95th — DESCENDING badge, -7.6%/1M                        | Frontier-relevant + paying scarcity prices for compute = margin pressure narrative. Long-horizon RS still elite; swing tape disagrees with the bull story right now. Patience entry, not chase |
| <b>MSFT / META</b> | “Disappointing” ASIC programs (named by Baker)                                                                                            | MSFT 26th/26th, META 23rd/25th — weak both horizons, -28% off highs | The episode quietly indicts both: heavy capex, no silicon edge, FCF down 80% as a group. Our RS data confirms the market agrees. Avoid until ASIC or model proof point                         |
| <b>AMZN</b>        | AWS origin story invoked as the EWS template; Bezos says orbital DCs ~6 yrs                                                               | 13th pct — FADING badge                                             | The historical analogy donor, not a recommendation. RS says distribution                                                                                                                       |
| <b>ORCL</b>        | Passed by SpaceX in the hyperscaler ranking (one line)                                                                                    | 54th / 41st, -44% off high                                          | The episode’s collateral damage — “got passed in 30 days.” Crowded 2025 AI-backlog story now mid-pack                                                                                          |

### D. Power, cooling & data-center BOM

| Name/Sector                             | What was said                                                                                                                                | Our data overlay                                           | How to think about it                                                                                                                                                                                                                                                                                                              |
|-----------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>GE Vernova (GEV)</b>                 | “If you’re Vernova with limited gas turbines — who do you sell to?” Speed premium accrues to suppliers                                       | +38.9% YTD but -15.4%/1M, -21% off high                    | Turbine scarcity is real but the stock already paid it. Baker’s “that game is over” applies                                                                                                                                                                                                                                        |
| <b>VRT (cooling) / CEG, VST (power)</b> | Not named directly; the “\$25B/GW ground BOM is inflationary” comment supports them long-term; orbital thesis undermines them long-long-term | VRT 16th pct DESCENDING -19%/1M; CEG -30% YTD; VST -9% YTD | The 2025 bottleneck trade is structurally supported but tactically broken. If orbital compute works (2028+), space bypasses the entire terrestrial BOM — these become the structural shorts of the orbital thesis. Two-stage view: hold off now (falling knife), revisit on consolidation, cap position life to pre-orbital window |
| <b>Behind-the-meter engineering</b>     | “Only 2-3 players can reliably engineer behind-the-meter DCs” (Gerstner, unnamed portfolio co.)                                              | —                                                          | Unnamed = likely an Altimeter private. Treat the claim as marketing until names surface                                                                                                                                                                                                                                            |

## E. Connectivity & telecom

| Name/Sector                                | What was said                                                                                                                                                     | How to think about it                                                                                                                                                                  |
|--------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Starlink (inside SpaceX)                   | <1% global household penetration; “hundreds of millions of terminals” possible; \$10B→\$50B by 2028 = only 0.3% of global telecom TAM; best+fastest+soon cheapest | The most “doable” lever per the table. Direct-to-cell is the expansion kicker (Gerstner’s Sand Hill dead-zone bit)                                                                     |
| Telecom incumbents (T, VZ, TMUS — unnamed) | Implicit: Starlink direct-to-cell takes share from the \$15T global telecom market                                                                                | The unspoken short side of the Starlink thesis. Satellite D2C at scale = MVNO-ization of rural/marine/air coverage first, urban later. Slow-burn structural pressure, not a 2026 trade |

## F. Software & internet (the bleeding sectors)

| Sector         | What was said                                                                               | Our data overlay                                            | How to think about it                                                                                                                                                                                                                                                                                    |
|----------------|---------------------------------------------------------------------------------------------|-------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Software (IGV) | “Down 8% on the year” (we measure –14%); Stripe’s 50M-line refactor as the existential demo | CRM –37%, NOW –33% YTD; NOW 82nd pct swing<br>RS = bouncing | The episode’s silent victim: Fable-5-class agents doing in a day what teams did in weeks is <i>deflationary for seat-based software</i> . The bounce in NOW is tradeable, the secular question is unanswered. Underweight stays correct until a software name shows <i>AI revenue</i> not <i>AI cost</i> |
| Internet (FDN) | “Down 16%”                                                                                  | FDN –1.5%, but mega-cap internet names –14 to –28%          | Same logic — agents disintermediate aggregators. Selective: GOOGL is also a frontier lab; pure aggregators aren’t                                                                                                                                                                                        |

## G. AI application layer (private, mentioned)

| Name                      | Context                                                         | Takeaway                                                                                                                                         |
|---------------------------|-----------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------|
| Harvey (legal AI)         | OSS + proprietary data + router beat Opus 4.7/4.8 at lower cost | The blueprint for app-layer economics: own data + routing = margin. Bearish for frontier-lab pricing power in verticals with data moats          |
| Replit (Amjad Masad cite) | “Coding may be the fastest path to AGI/ ASI”                    | Context for why coding assets (Cursor) carry strategic premiums                                                                                  |
| Stripe                    | 50M-line Ruby refactor in 1 day w/ Fable 5                      | The enterprise-adoption proof point of the episode                                                                                               |
| Fireworks                 | RL/fine-tuning infra used by Harvey                             | Picks-and-shovels of the routing era                                                                                                             |
| Kimi / Moonshot (K2.5)    | Base model for Cursor’s Composer                                | Chinese OSS base models are inside US production stacks — JPM’s “loathe to use Chinese models” comment cuts against this; US OSS lab gap is real |

## 8) Companion Files

- [...\\_claims.md](#) — exhaustive claim-by-claim table with timestamps (every number in the episode)
- [...\\_participants.md](#) — sourced research brief on all four participants, firm stakes, and externally-confirmed IPO/ deal record (37 sources)
- [transcripts/20260611\\_bg2\\_spacex\\_ipo\\_fable5\\_ai\\_capex.{json,txt}](#) — raw archived transcript + description

## 9) PM Bottom Line

---

**Action → Monitor (SpaceX IPO), Rotate-watch (NVDA/AVGO vs bottleneck winners), Hold underweight software.**

The single most useful sentence in the episode is Baker's: *"finding the next bottleneck was the last game — that game is over."* Combined with verified price action (MU/MRVL/AMD +128-249% YTD vs NVDA +10%, AVGO +11%), the table is signaling rotation **from scarcity winners back to the platform names** — while simultaneously trimming gross exposure into a seasonal consolidation they all expect. On SpaceX itself: the analysis is better than the average IPO pitch, but it IS a pitch — from holders, 48h before pricing, with the bear chair empty. Their own >50%-drawdown chart plus the "base position + ballast" sizing doctrine is the actionable takeaway, not "set it and forget it."

**Key watches:** 1. **Starship second-stage recovery attempt (H2 2026)** — binary gate for Starlink V3 + orbital math (T2). 2. **First Grok 4.3 (1.5T) + Cursor-data model release** — the highest-information event for xAI's value (T4). 3. **Anthropic revenue disclosures** — the load-bearing wall of the entire 2026 AI tape (their words). 4. **Token-mix data (silicon data index)** — open-source share shift + summer seasonality = the early warning on token pricing (T6/T8, Clark's caveat). 5. **SpaceX post-IPO tape** — lockup dribble vs index-inclusion flows; their own chart says >50% max drawdown is the base rate. 6. **Vera Rubin allocations** — confirm/deny Clark's "xAI got 20%" channel check. 7. **Macro spoilers** — CPI >4%, oil ~\$100, Iran: the consolidation catalyst they're positioned for and retail isn't.

---

*Analysis: Apexium research. Transcript + description archived in [transcripts/](#). Price/RS overlays from Apexium local bars + RS dashboard snapshot 2026-06-11. Claims inventory by extraction agent, verified spot-wise against transcript.*